Bootcamp 6: Reinforcement Learning

William H. Guss, James Bartlett
{wguss, james}@ml.berkeley.edu
Machine Learning at Berkeley

April 22, 2016

Overview

Problem: ML for Pacman.

How would you solve pacman with machine learning?

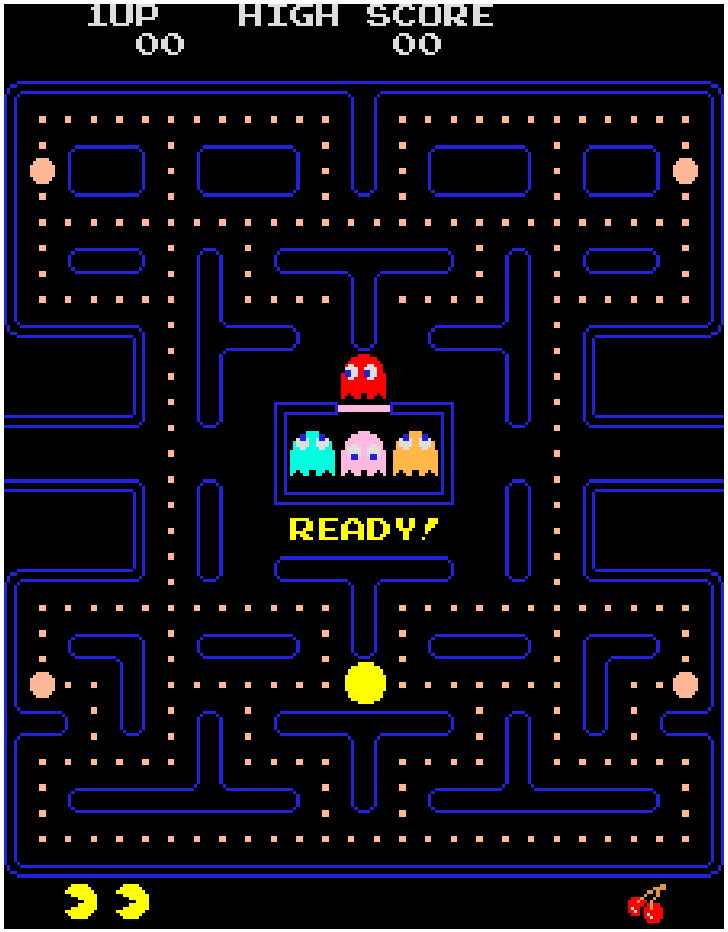

Problem: ML for Pacman.

How would you solve pacman with machine learning?

Find a model which takes screen pixels to actions:

$$\pi_\theta : s_t \mapsto a_t.$$

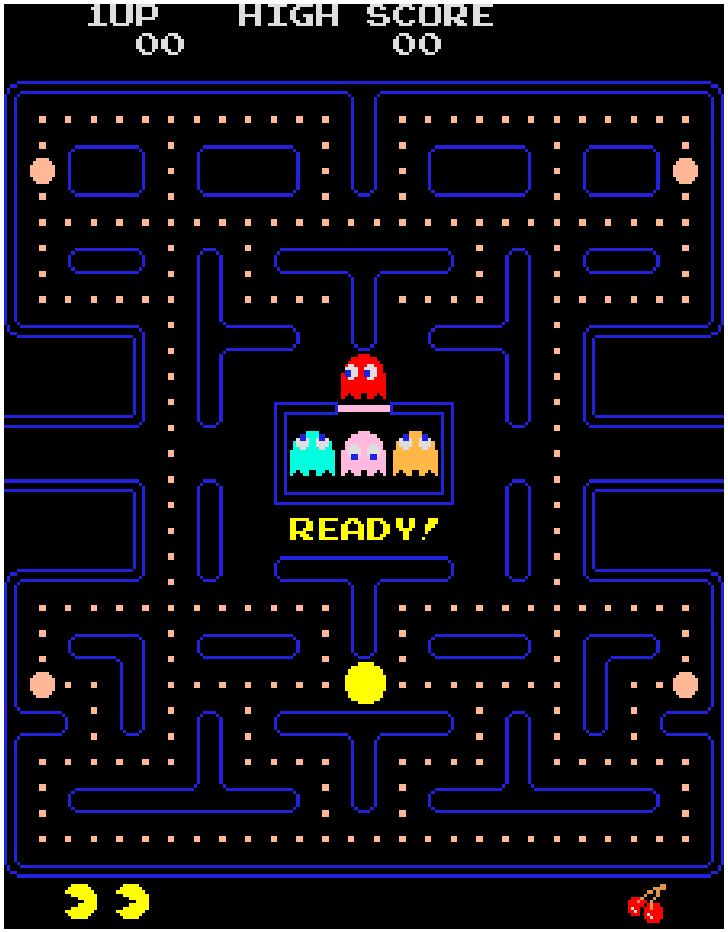

Problem: ML for Pacman.

How would you solve pacman with machine learning?

Find a model which takes screen pixels to actions:

$$\pi_\theta : s_t \mapsto a_t.$$

What is your loss function? Data?

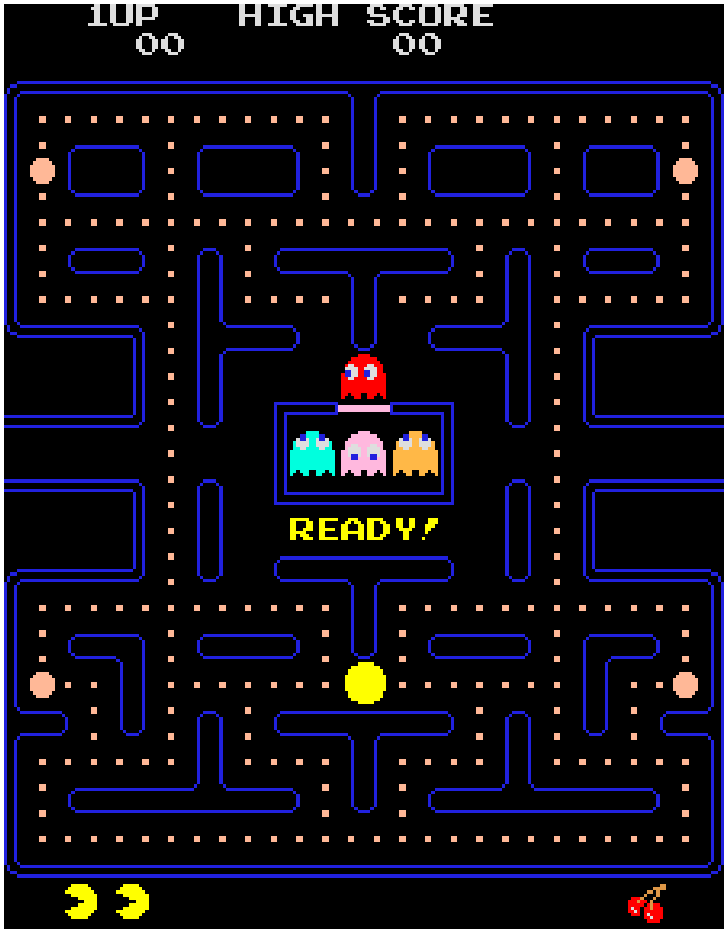

Problem: ML for Pacman.

$\Longrightarrow$

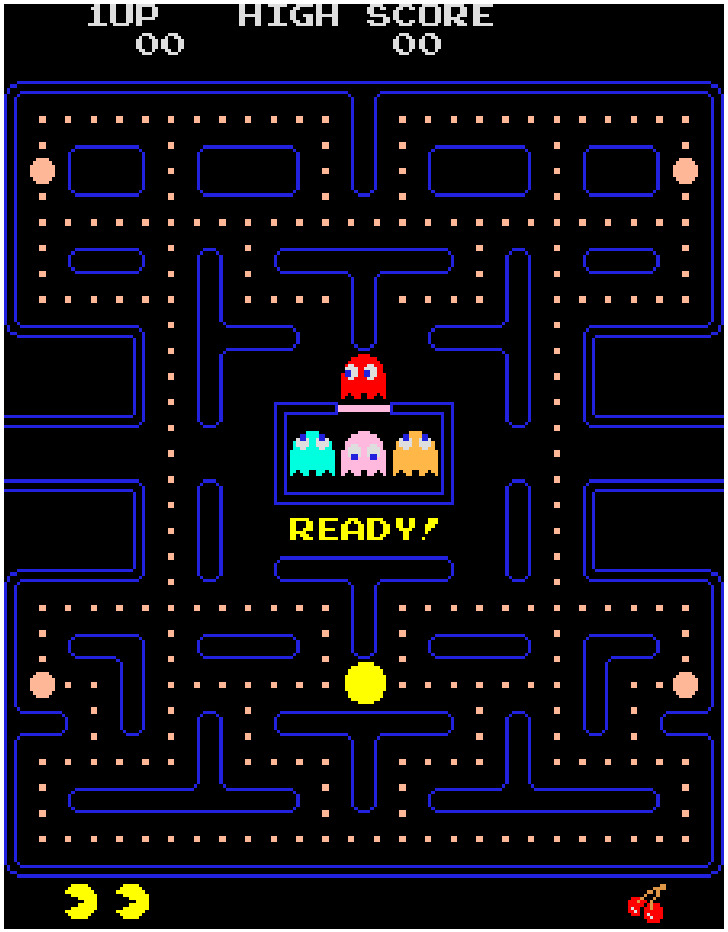

Solution: Reinforcement Learning

- Supervised learning is *not* the most general formulation of learning.

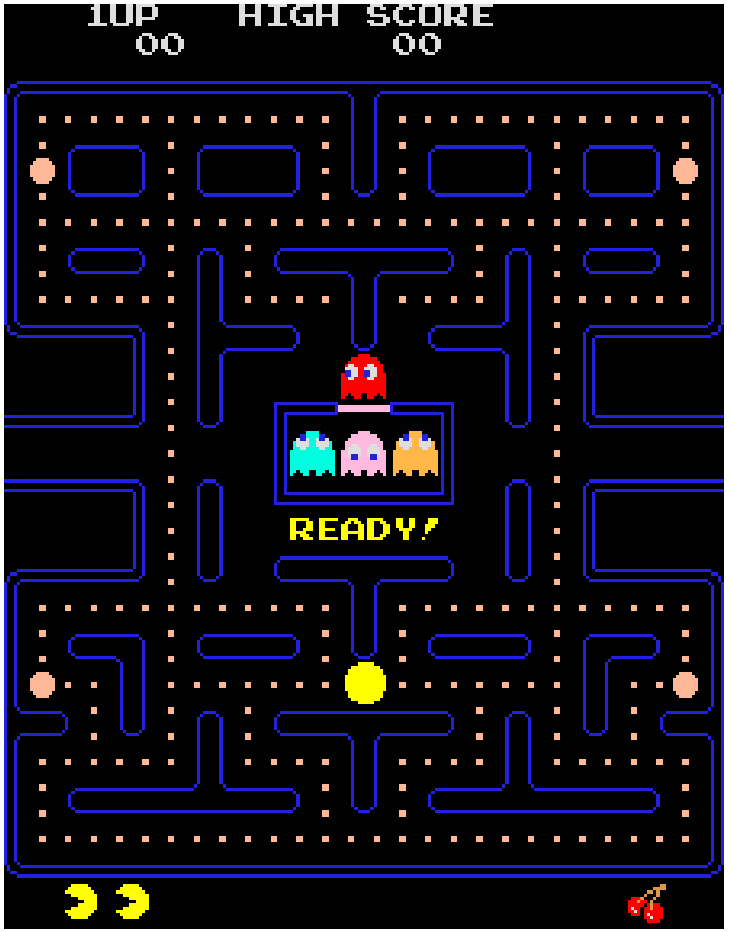

Solution: Reinforcement Learning

- Supervised learning is *not* the most general formulation of learning.
- Humans learn through reward and penalty

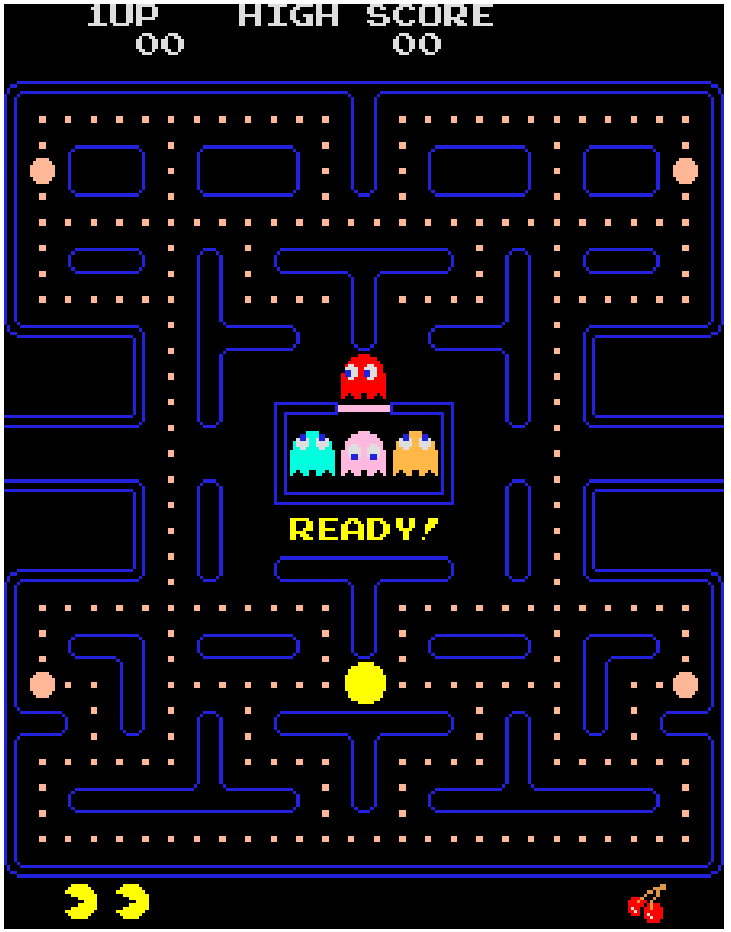

Solution: Reinforcement Learning

- Can we make algorithms which improve with crude reward signals?

Machine learning without explicit objective functions

$\Downarrow$

Reinforcement Learning (RL)

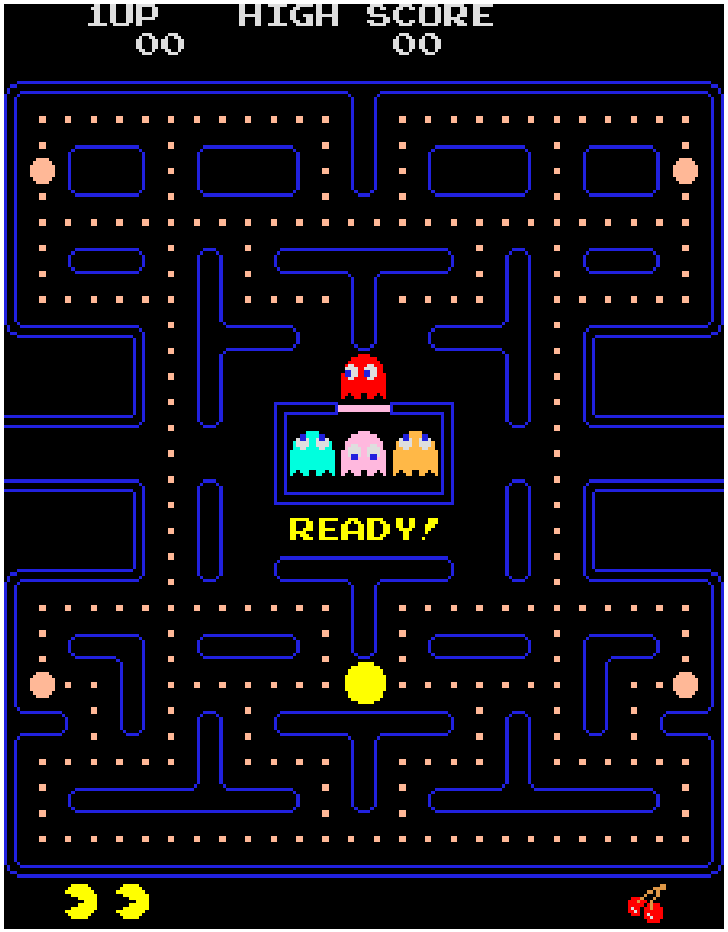

The Core Idea

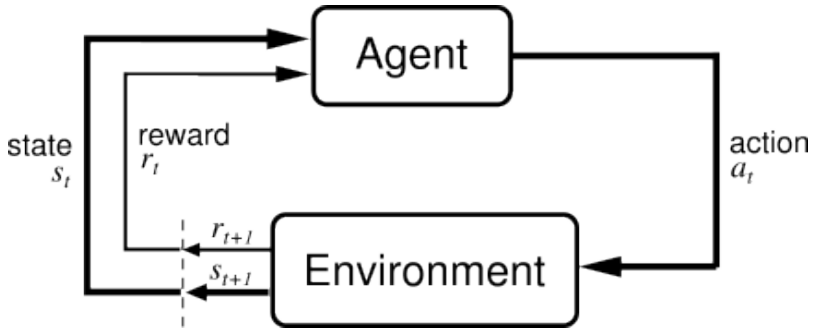

- Models (agents) take action a_t in some environment.
- Environment provides state s_t, reward r_t.
- Models learn to maximize reward r_t, $\forall t$.

Markov Decision Process (MDP)

Environment, $E = (\mathcal{S}, \mathcal{A}, \mathcal{R}, \rho, r)$.

1. State space, $\mathcal{S}$
2. Action space, $\mathcal{A}$
3. Reward space, $\mathcal{R}$
4. Transition distribution, $\rho(s' \mid s, a)$. Given a previous state s and action a, environment gives s'.
5. Reward function $r(s, a) \in \mathcal{R}$.

Markov Property: $\rho(s' \mid s, a)$ depends only on s, a not previous states!

Markov Decision Process (MDP)

ML@B

Example MDP

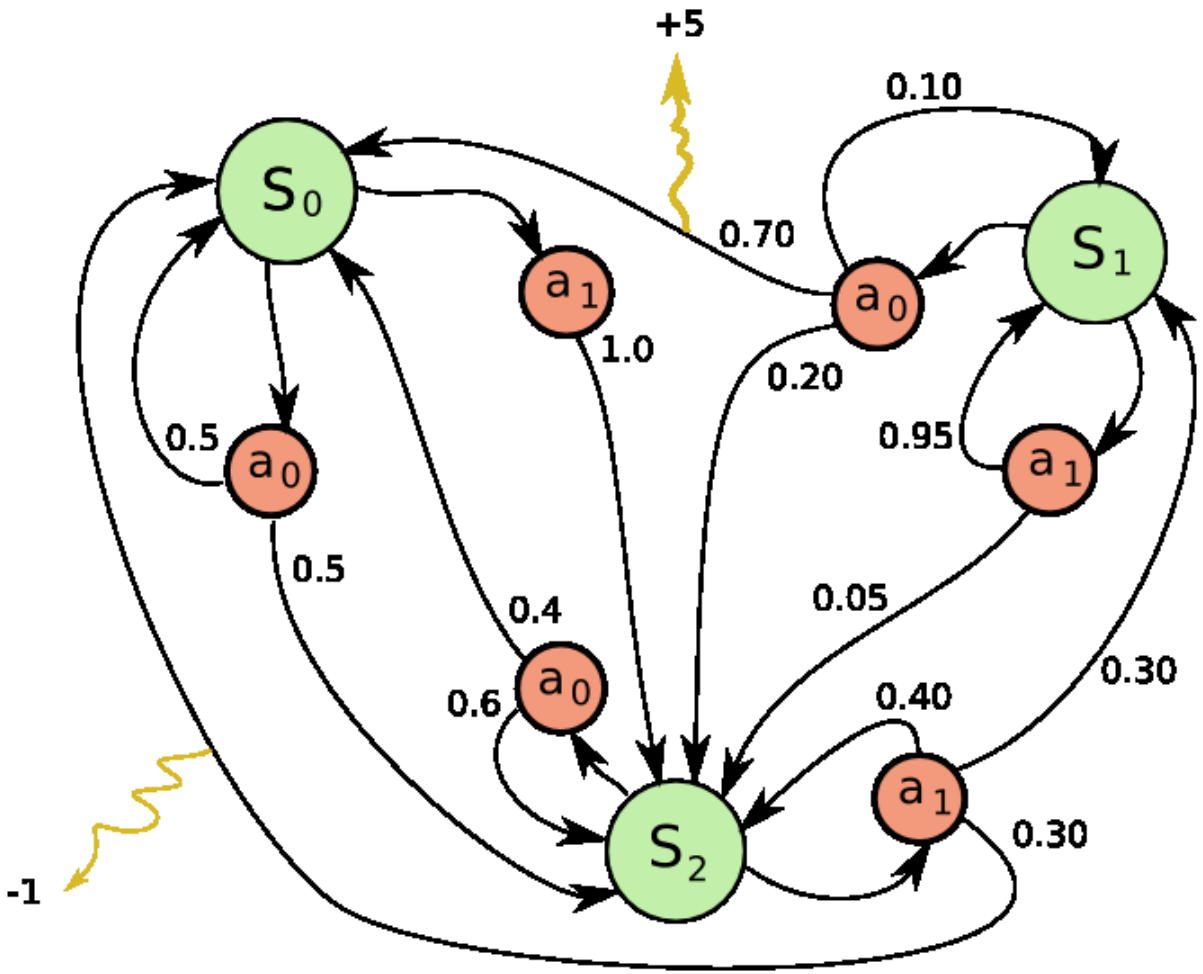

Pacman as an MDP

ML@B

- $\mathcal{S} = \mathbb{R}^{256 \times 256}$, images as state space.
- $\mathcal{A} = \{\uparrow, \downarrow, \rightarrow, \leftarrow\}$, joystick as action space.
- $r(s_t, a_t) =$ change in score.
- $\rho(s_{t+1} \mid s_t, a_t) =$ next frame of game after joystick action a_t.

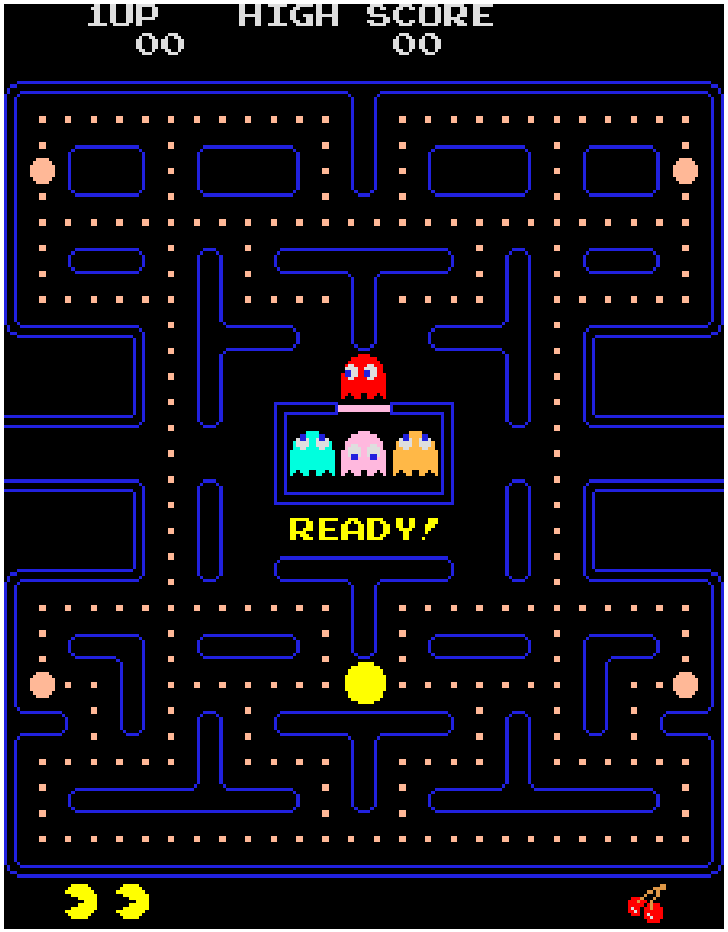

Policies/Agents

ML@B

Two different types of agents

- Deterministic policy $a = \pi(s)$ acts in E.
- Stochastic policy $a \sim \pi(a|s)$ gives a probability distibution over actions.

Policy Trajectories

$$s_1 \xrightarrow{\pi} a_1 \xrightarrow{\rho,r} s_2, r_2 \xrightarrow{\pi} a_2 \xrightarrow{\rho,r} \cdots$$

Value under a policy

The **state value** is a function of a given state for an agent π defined as

$$V^{\pi}(s_t) = \mathbb{E}\left[\sum_{n=t+1}^{\infty} \gamma^n r(s_n, \pi(s_n))\right]$$

1. γ is the discount factor
2. $\pi(s_n)$ is the action the agent π makes after seeing state s_n.
3. $r(s_n, \pi(s_n))$ is the reward the agent gets from taking that action.

Value under a policy

The **state value** is a function of a given state for an agent π defined as

$$V^{\pi}(s_t) = \mathbb{E}\left[\sum_{n=t+1}^{\infty} \gamma^n r(s_n, \pi(s_n))\right]$$

$$s_t \xrightarrow{\pi} a_1 \xrightarrow{\rho,r} s_2, r_2 \xrightarrow{\pi} a_2 \xrightarrow{\rho,r} \cdots$$

$$s_t \xrightarrow{\pi} a_0 \xrightarrow{\rho,r} s_7, r_7 \xrightarrow{\pi} a_3 \xrightarrow{\rho,r} \cdots$$

Value under a policy

The **state-action value** for an agent π is defined such that

$$Q^{\pi}(s_t, a_t) = \mathbb{E}\left[\underbrace{r(s_t, a_t)}_{\text{reward for } a_t} + V^{\pi}(s_t)\right]$$

- Given some state s_t, the *best* agent, π^* is one that take action

$$a_t = \underset{a}{\operatorname{argmax}}\, Q(s_t, a).$$

Problems in Reinforcement Learning

- **Policy Optimization:** maximize the expected reward with respect to a policy π;

$$\pi^* = \underset{\pi}{\operatorname{argmax}}\, \mathbb{E}\left[\sum_{t=0}^{\infty} r_t\right]$$

- **Policy Evaluation:** Given some fixed policy π compute expected return.
  - Computing Q^π, V^π, and other expectations on policy rollout.
  - Lets us perform policy optimization!

Assorted Algorithms

We'll go over:

- Behavioral Cloning
- Q-Learning
- Policy Iteration

Learn at home:

- Value iteration
- Temporal Difference Methods
- Inverse Reinforcement Learning.

Behavioral Cloning

ML@B

Behavioral Cloning: Supervised learning in MDPs using and expert agent expert π^*!

Behavioral Cloning

ML@B

Behavioral Cloning: Supervised learning in MDPs using and expert agent expert π^*!

Given expert examples $\mathcal{D} = (s_t, a_t = \pi^*(s_t))$ and a model π_θ find θ^* st

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \mathcal{L}(a_t, \pi_\theta(s_t)).$$

where $\mathcal{L}$ is some loss function.

Behavioral Cloning

ML@B

Behavioral Cloning: Supervised learning in MDPs using and expert agent expert π^*!

Given expert examples $\mathcal{D} = (s_t, a_t = \pi^*(s_t))$ and a model π_θ find θ^* st

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \mathcal{L}(a_t, \pi_\theta(s_t)).$$

where $\mathcal{L}$ is some loss function.

- Show, don't tell!

Behavioral Cloning

Behavioral Cloning: Supervised learning in MDPs using and expert agent expert π^*!

Given expert examples $\mathcal{D} = (s_t, a_t = \pi^*(s_t))$ and a model π_θ find θ^* st

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \mathcal{L}(a_t, \pi_\theta(s_t)).$$

where $\mathcal{L}$ is some loss function.

- Show, don't tell!
- No complicated machinery, just standard ML.

Behavioral Cloning

ML@B

Issue: Compounding Error

Given some irreducible error $\epsilon = 0.001$

- $\mathcal{L}(a_0, \pi_\theta(s_0)) = \epsilon$
- $\mathcal{L}(a_1, \pi_\theta(s_1)) = 2\epsilon$
- $\mathcal{L}(a_2, \pi_\theta(s_2)) = 3\epsilon$
- $\mathcal{L}(a_3, \pi_\theta(s_3)) = 4\epsilon$
- $\mathcal{L}(a_4, \pi_\theta(s_4)) = 5\epsilon$

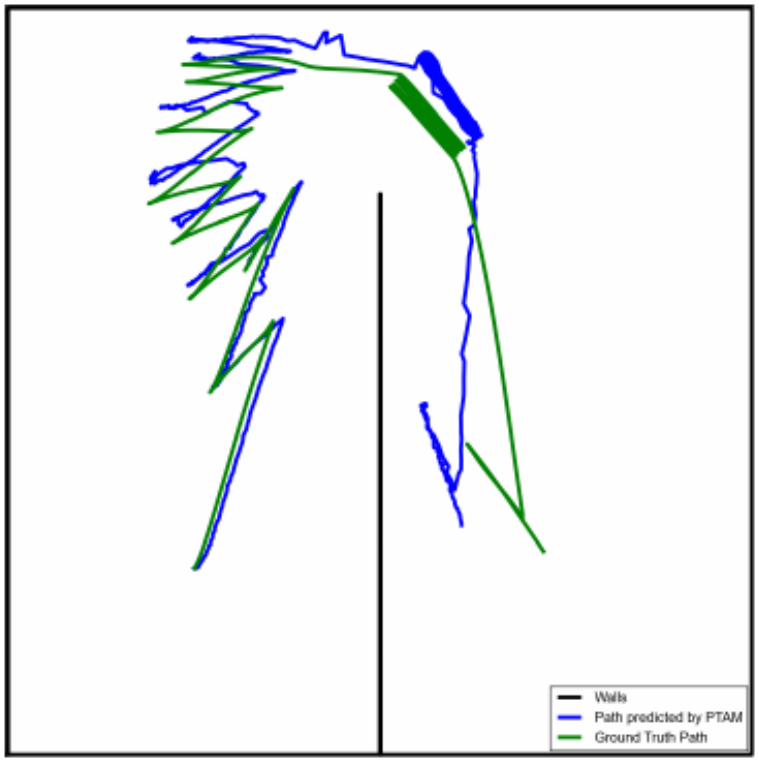

Behavioral Cloning

ML@B

Issue: Distribution Mismatch

- States expert dataset $\mathcal{D}$ generated by π^* have different distribution than those generated by π_θ.

  $\implies$ No self correction.

Behavioral Cloning

Issue: Distribution Mismatch

- States expert dataset $\mathcal{D}$ generated by π^* have different distribution than those generated by π_θ.
  $\implies$ No self correction.

Solution: DAgger.

- Do BC on $\mathcal{D}$ and generate E_0 states generated by π_θ.
- Label E_0 with expert level actions and add to $\mathcal{D}$.

Q-Learning (State-action Value Iteration)

Goals of Q-learning

1. Approximate Q^{π^*}, the Q function of the optimal agent, as $Q(s_t, a_t)$.

Q-Learning (State-action Value Iteration)

Goals of Q-learning

1. Approximate Q^{π^*}, the Q function of the optimal agent, as $Q(s_t, a_t)$.
2. Using Q, find the agent, π, that best approximates the optimal agent, π^*.

Q-Learning (State-action Value Iteration)

How do we define best?

Q-Learning (State-action Value Iteration)

How do we define best?

Given some state s_t, the **best** agent, π^* is one that takes action

$$a_t = \arg\max_a Q(s_t, a).$$

Q-Learning (State-action Value Iteration)

An example: Frozen Lake Problem

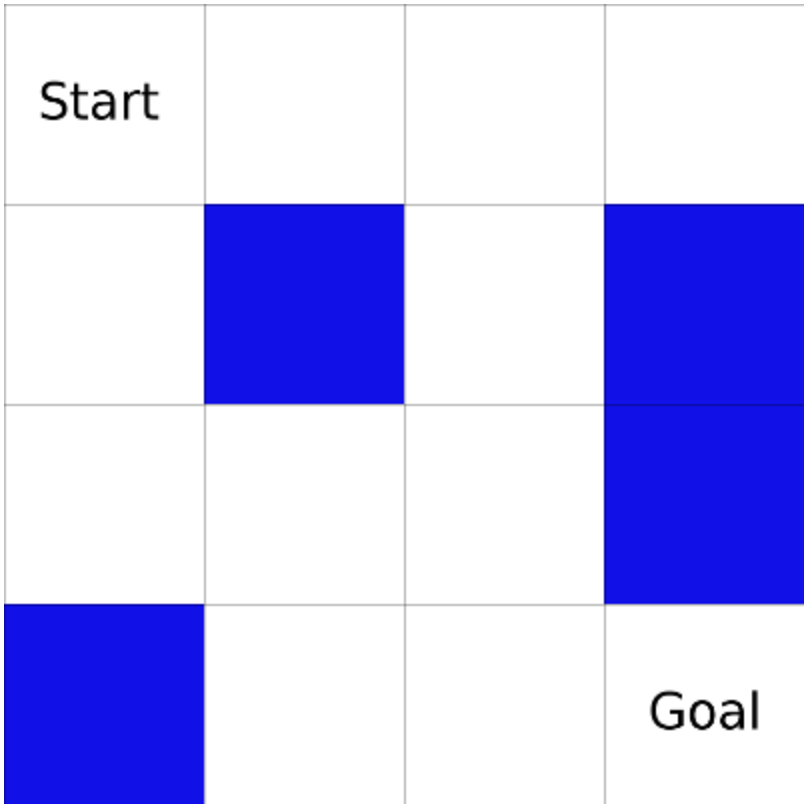

Q-Learning (State-action Value Iteration)

- 100 reward for reaching the goal
- 0 otherwise

How do we keep track of this long term reward?

Q-Learning (State-action Value Iteration)

- 100 reward for reaching the goal
- 0 otherwise

How do we keep track of this long term reward?

Q function

Q-Learning (State-action Value Iteration)

How do we actually calculate the Q function?

Q-Learning (State-action Value Iteration)

How do we actually calculate the Q function?

The Bellman Equation.

Q-Learning (State-action Value Iteration)

How do we actually calculate the Q function?

The Bellman Equation.

$$Q^{\pi}(s_t, a_t) = r_t + \gamma Q^{\pi}(s_{t+1}, \pi(s_{t+1}))$$

Q-Learning (State-action Value Iteration)

One Q-Learning Algorithm: Tabular Q-Learning

- Explore the environment
- On the way, use the Bellman equation to store a table of expected future reward (Q) for each state-action pair.
- Use this table to pick the best possible action for any given state.

Q-Learning (State-action Value Iteration)

An example update for Frozen Lake.

Suppose our stored Q table looks like so:

Up	Down	Left	Right
0	65	0	40
0	0	0	0
0	0	0	0
0	0	0	0
50	75	30	20
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Q-Learning (State-action Value Iteration)

ML@B

An example update for Frozen Lake.

Then suppose our agent moves **down** from the starting square

Q-Learning (State-action Value Iteration)

ML@B

An example update for Frozen Lake.

Then we update using the Bellman equation.

$$Q(s_{t+1}, a_{t+1}) = Q(s_t, a_t) + \alpha(r_t + \gamma(\max_a Q(s_t, a) - Q(s_t, a_t))$$

Up	Down	Left	Right
0	65	0	40
0	0	0	0
0	0	0	0
0	0	0	0
50	75	30	20
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Q-Learning (State-action Value Iteration)

An example update for Frozen Lake.

The table now looks like so:

Up	Down	Left	Right
0	70	0	40
0	0	0	0
0	0	0	0
0	0	0	0
50	75	30	20
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

Policy Iteration

Policy Iteration: Given access to the MDP, use policy evaluation to iteratively serach for better policies!

- Choose a policy at random, π.

Policy Iteration

Policy Iteration: Given access to the MDP, use policy evaluation to iteratively serach for better policies!

- Choose a policy at random, π.
- Alternate between
  - Evaluate policy $\pi \rightarrow V^{\pi}$.

Policy Iteration

ML@B

Policy Iteration: Given access to the MDP, use policy evaluation to iteratively serach for better policies!

- Choose a policy at random, π.
- Alternate between
  - Evaluate policy $\pi \rightarrow V^{\pi}$.
  - Set new policy to be greedy policy for V^{π}

$$\begin{aligned}\pi(s) &:= \underset{a}{\operatorname{argmax}}\, \mathbb{E}\left[R(s,a) + \gamma V^{\pi}(s')\right] \\ &:= \underset{a}{\operatorname{argmax}}\, Q^{\pi}(s,a)\end{aligned}$$

Policy Iteration

Policy Iteration: Given access to the MDP, use policy evaluation to iteratively serach for better policies!

- Choose a policy at random, π.
- Alternate between
  - Evaluate policy $\pi \to V^\pi$.
  - Set new policy to be greedy policy for V^π

$$\begin{aligned}\pi(s) &:= \operatorname*{argmax}_a \mathbb{E}\left[R(s,a) + \gamma V^\pi(s')\right] \\ &:= \operatorname*{argmax}_a Q^\pi(s,a)\end{aligned}$$

Note: Learn Q^π using Q-learning without argmax.

$$Q^\pi(s_t, a_t) = r(s_t, a_t) + \gamma Q^\pi(s_{t+1}, \pi(s_{t+1})$$

Policy Iteration

ML@B

Example: Deep Determisitic Policy Gradient

1. Actor neural network $\pi_\theta : \mathcal{S} \to \mathcal{A}$
2. Critic network $Q^\pi : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$
3. Performance of π is $Q^\pi(s_t, \pi(s_t))$. **Maximize performance!** $\nabla_\theta Q^\pi(s_t, a_t) = \nabla_a Q^\pi(s_t, a) \cdot \nabla_\theta \pi_\theta(s_t)$

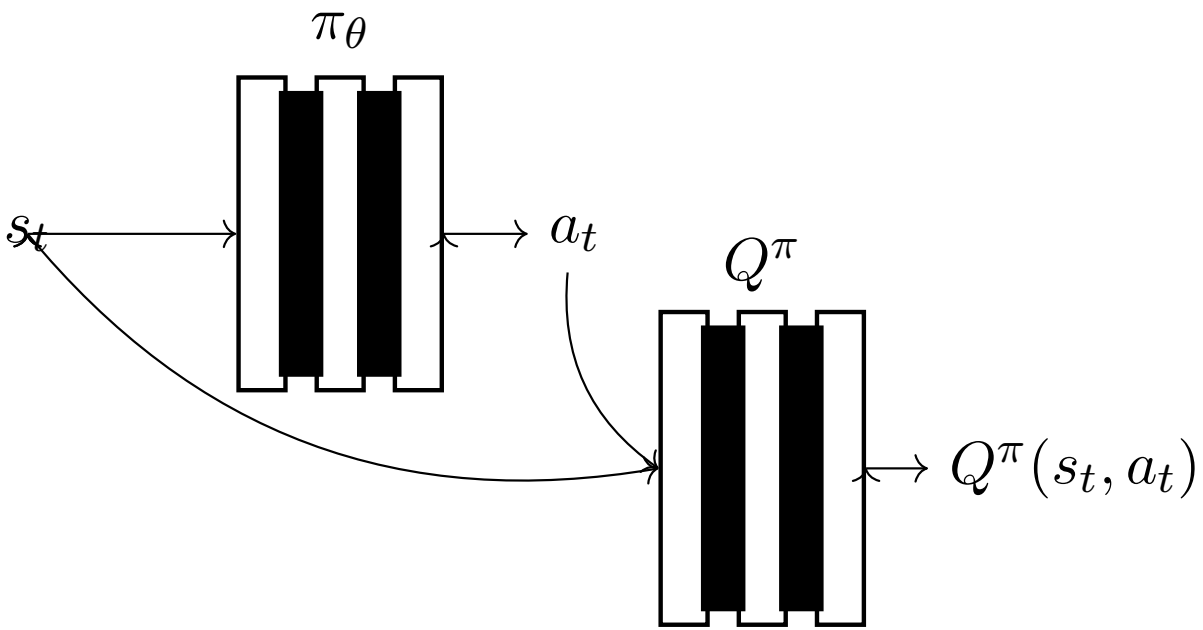

and many more...

Papers and links to descriptions of other algorithms will be posted on github.com/mlberkeley/bootcamp.

Practical Reinforcement Learning

Process

1. Clearly define the environment.
   - What is state space, $\mathcal{S}$? Action space $\mathcal{A}$? Low dimensions is better!
   - Use tools like OpenAI Gym to make a simulator.
2. Choose a good reward function.
   - Don't make sparse reward functions:

     +1 for winning vs. guiding agent to win.

   - Penalize for time constraints!
   - Break the problem in to smaller goals with rewards $r_1, r_2, \ldots$. Then $r(s,a) = \sum_i w_i r_i(s,a)$.

Process

ML@B

3. Choose a model.
   - Deep reinforcement learning is king!
     - Discrete $\mathcal{A} \rightarrow DQN, DDQN$
     - Continuous $\mathcal{A} \rightarrow DDPG, NAF, A3C, TRPO.$
   - Less hyperparameters better.
   - Don't rule models out too early.
4. Train the model against the environment.
   - Choose a good exploration strategy.
   - RL can take a long time (days - months), let models train to completion.
   - Do a sweep on hyperparameters.
   - Overfitting (to optimal policy) is good.
   - *(You will fail here many times, go to step 3.)*
5. Profit

ProTips^{tm}

ML@B

- **Fastest way to learn RL:** Implement landmark papers in tensorflow, Atari DQN, DPG, ...
- **Write unit tests before you train!**
- **Use an autodiff framework:** Tensorflow, PyTorch, ...
- **Avoid pixel data at all costs:** CNNs/CV algorithms make training very slow!
  - Test RL models on featurized data in simualtion, then move to pixel based models for real data.

Good luck!

Questions?